

# Julian Skirzyński

CURRICULUM VITAE — MAY 2026

[jskirzynski@ucsd.edu](mailto:jskirzynski@ucsd.edu)

[www.jskirzynski.com](http://www.jskirzynski.com)

|                    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |                                            |
|--------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------|
| EDUCATION          | <b>University of California, San Diego</b><br><i>Ph.D. Candidate in Computer Science &amp; Engineering</i><br>Thesis: Foundations of Human AI Interaction<br>Advisor: Berk Ustun                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        | 2022 – PRESENT                             |
|                    | <b>McGill University</b><br><i>M.S. in Computer Science</i><br>Thesis: <a href="#">Language-Conditional Imitation Learning</a><br>Advisor: David Meger                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  | 2017 – 2020                                |
|                    | <b>University of Warsaw</b><br><i>M.S. in Cognitive Science</i><br><i>B.S. in Mathematics, Cognitive Science</i><br>Advisors: Andrzej Skowron; Piotr Wasilewski                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         | 2012 – 2018                                |
| ACADEMIC POSITIONS | <b>Max Planck Institute for Intelligent Systems, Germany</b><br><i>Research Scientist</i><br>Projects: Interpretable RL Policies, Improving Human Planning, Discovering Human Planning Strategies<br>Advisor: Falk Lieder                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               | 2019 – 2023                                |
| RESEARCH INTERESTS | <b>Areas:</b> Machine Learning, Cognitive Science, Human-Computer Interaction<br><b>Topics:</b> Decision-Making, Interpretability, AI Alignment, Reinforcement Learning, Experimental Design<br><b>Applications:</b> Social Sciences, Medicine, Consumer Finance, Criminal Justice                                                                                                                                                                                                                                                                                                                                                                                                                                      |                                            |
| AWARDS & HONORS    | Pierre Arbour Foundation Scholarship<br>McGill University Graduate Excellence Award<br>McGill - University of Warsaw Exchange Scholarship<br>University of Warsaw Academic Excellence Scholarship                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       | 2018 – 2019<br>2018<br>2015<br>2014 – 2017 |
| PREPRINTS          | <ol style="list-style-type: none"><li><a href="#">What Went Wrong? Structured Self-Debugging for LLM Agents</a><br/>Julian Skirzyński, Berk Ustun<br/><i>In Submission, 2026</i></li><li><a href="#">The Illusion of Grounded Reasoning in LLMs</a><br/>Julian Skirzyński, Berk Ustun<br/><i>In Submission, 2026</i></li><li><a href="#">Measuring What Matters: Synthetic Benchmarks for Concept Bottleneck Models</a><br/>Julian Skirzyński, Harry Cheon, Shreyas Kadekodi, Meredith Stewart, Berk Ustun<br/><i>In Submission, 2026</i></li><li><a href="#">On the Value of Interpretability in Human Decision-Making</a><br/>Julian Skirzyński, Elena Glassman, Berk Ustun<br/><i>In Preparation, 2025</i></li></ol> |                                            |
| PAPERS             | <ol style="list-style-type: none"><li><a href="#">Quantifying Cognitive Bias Induction in LLM-Generated Content</a><br/>Abeer Alessa, Akshaya Lakshminarasimhan, Param Somane, Julian Skirzyński, Julian McAuley, Jessica Echtermoff<br/><i>International Joint Conference on Natural Language Processing, 2025</i></li></ol>                                                                                                                                                                                                                                                                                                                                                                                           |                                            |

\*EQUAL CONTRIBUTION



4. [Discrimination Exposed? On the Reliability of Explanations for Discrimination Detection](#)  
Julian Skirzyński, David Danks, Berk Ustun  
*ACM Conference on Fairness, Accountability, and Transparency*, 2025
5. [Automatic Discovery and Description of Human Planning Strategies](#)  
Julian Skirzyński, Yash Raj Jain, Falk Lieder  
*Behavior Research Methods*, 2023
6. [Boosting Human Decision-making with AI-Generated Decision Aids](#)  
Frederic Becker\*, Julian Skirzyński\*, Bas van Opheusden, Falk Lieder  
*Computational Brain & Behavior*, 2022
7. [Automatic Discovery of Interpretable Planning Strategies](#)  
Julian Skirzyński, Frederic Becker, Falk Lieder  
*Machine Learning*, 2021
8. [Object \[Re\] Cognition with Similarity](#)  
Łukasz Sosnowski, Julian Skirzyński  
*International Conference on Information Processing and Management of Uncertainty in Knowledge-Based Systems*, 2018
9. [A Framework for Analysis of Granular Neural Networks](#)  
Julian Skirzyński  
*International Joint Conference on Rough Sets*, 2017

REFEREED  
WORKSHOP  
PAPERS

9. [On Interpretability and Overreliance](#)  
Julian Skirzyński, Elena Glassman, Berk Ustun  
*Interpretable AI: Past, Present and Future, NeurIPS Workshop*, 2024
10. [Language-Conditional Imitation Learning](#)  
Julian Skirzyński, Bobak Baghi, David Meger  
*Visually Grounded Interaction and Language, NAACL Workshop*, 2021

TEACHING

**UCSD Computer Science & Engineering Department** 2026  
[CSE101 – Design & Analysis of Algorithms](#)  
*Teaching Assistant*  
 Led discussion sessions for 300+ undergraduate students on the basics of theoretical computer science (graph theory, runtime, NP-completeness, etc.). Held weekly office hours and corrected homework.

**UCSD Halicioğlu Data Science Institute** 2023  
[DSC291 – Interpretability & Explainability in Machine Learning](#)  
*Guest Lecturer & Teaching Assistant*  
 Co-designed curriculum and held weekly office hours for serving 20+ PhD/MS students. Delivered guest lectures on ML interpretability methods and cognitive biases in AI-assisted decision-making. Completed teaching development workshop on graduate-level instruction.

SOFTWARE  
 [GitHub](#)

[Strategy Extraction from RL Policies](#) – Algorithm to extract interpretable decision trees from RL policies  
[Human Planning Strategy Analysis](#) – Framework for identifying strategies used in human planning tasks

SELECTED  
INDUSTRY  
POSITIONS

**Educational Entertainment One, Warsaw, Poland** 2021 – 2024  
*Lead Technical Architect*  
 Designed algorithms (AI, NLP) and supported the production process for a story-driven mobile game for learning English.

|                     |                                                                                                    |                |
|---------------------|----------------------------------------------------------------------------------------------------|----------------|
| ACADEMIC<br>SERVICE | JOURNAL REVIEWING                                                                                  |                |
|                     | Machine Learning                                                                                   | 2022           |
|                     | CONFERENCE PROGRAM COMMITTEE                                                                       |                |
|                     | NeurIPS – Conference on Neural Information Processing Systems                                      | 2023 – PRESENT |
|                     | ICML – International Conference on Machine Learning                                                | 2025 – PRESENT |
|                     | ICLR – International Conference on Learning Representations                                        | 2024 – PRESENT |
|                     | AAAI Conference on Artificial Intelligence                                                         | 2025 – PRESENT |
|                     | FAccT – ACM Conference on Fairness, Accountability and Transparency                                | 2022 – PRESENT |
|                     | ICML Workshop RL4RealLife – International Conference on Machine Learning                           | 2021           |
|                     | IPMU – Information Processing and Management of Uncertainty in Knowledge-Based Systems             | 2018           |
| PERSONAL            | Language Skills : English, Polish, German (Conversational)                                         |                |
|                     | Software Skills : Python, R, C++, Flask, AWS, PyTorch, CPLEX, JavaScript, Jira                     |                |
|                     | Interests : Soccer, Groundhopping, Traveling, Fantasy Literature, Record Collecting                |                |
|                     | Other : Peer tutoring, Co-author of “Triozy polskie”, a textbook for learning Polish by foreigners |                |